

Erarica Mehra
NUID: 002113683

Assignment 4:
Parallel Sorting

Tasks

Implemented **Parallel Sorting** such that large arrays can be sorted based on the following schemes:

- Determine an optimal value for **cutoff**. If elements in the array are less than cutoff, then switch to system sort.
- Modify the cutoff values while sorting the array partitions asynchronously.
- Determining **recursion deception** (that is the number of threads). These threads sort the array parallelly.
- Deciding an ideal number of threads to efficiently sort the array. Partitioning must be stopped if recursion depth reaches $\log n$.
- Combination of these schemes: Here I am taking an initial thread count(2) and an initial array size value(1000000) from the command line and doubling these on each iteration, while monitoring the cutoff values and performance of parallel sorting of a large array.

Relationship/Conclusion:

- Parallel Sorting is dependent on balancing the load among the different processors of the system. Efficiency of the parallel sort is highly dependent on the cutoff values. The cutoff value determines when to switch to system sort.
- When the cutoff values are as high as the size of the array and the number of threads are high (that is 32), we can switch to system sort.
- When the cutoff is low and the number of threads are high, the time taken to parallel sort the array increases. Hence the performance of the parallel sort decreases.
- Therefore we can conclude that when the cutoff values are low, the number of threads should be low for better performance. Therefore, it can be rightly said that cutoff values are inversely proportional to the number of threads. Based on the given trials and observations we can decide that the number of threads in this

case can be 4. With 4 threads, the system is efficiently using the processors of the system and the overhead to switch between threads/tasks is minimum.

Observations/ Evidence :

All of the given data is captured in respective CSVs and uploaded with the project in the results folder.

1. Screenshot of the console output:

I have taken a different array sizes with different degrees of parallelization that has produced the following output.

```
Enter initial thread count :  
2  
Enter initial array size :  
1000000  
Degree of parallelism: 2  
Size of the array: 1000000  
cutoff: 510000    10 times Time :832ms  
cutoff: 520000    10 times Time :448ms  
cutoff: 530000    10 times Time :441ms  
cutoff: 540000    10 times Time :457ms  
cutoff: 550000    10 times Time :464ms  
cutoff: 560000    10 times Time :451ms  
cutoff: 570000    10 times Time :439ms  
cutoff: 580000    10 times Time :446ms  
cutoff: 590000    10 times Time :453ms  
cutoff: 600000    10 times Time :455ms  
cutoff: 610000    10 times Time :458ms  
cutoff: 620000    10 times Time :455ms  
cutoff: 630000    10 times Time :440ms  
cutoff: 640000    10 times Time :446ms  
cutoff: 650000    10 times Time :446ms  
cutoff: 660000    10 times Time :444ms  
cutoff: 670000    10 times Time :466ms  
cutoff: 680000    10 times Time :445ms  
cutoff: 690000    10 times Time :498ms  
cutoff: 700000    10 times Time :451ms  
cutoff: 710000    10 times Time :447ms  
cutoff: 720000    10 times Time :462ms  
cutoff: 730000    10 times Time :458ms  
cutoff: 740000    10 times Time :477ms  
cutoff: 750000    10 times Time :439ms  
cutoff: 760000    10 times Time :442ms  
cutoff: 770000    10 times Time :448ms  
cutoff: 780000    10 times Time :439ms  
cutoff: 790000    10 times Time :439ms  
cutoff: 800000    10 times Time :439ms  
cutoff: 810000    10 times Time :440ms  
cutoff: 820000    10 times Time :455ms  
cutoff: 830000    10 times Time :466ms  
cutoff: 840000    10 times Time :444ms
```

```
Degree of parallelism: 4
Size of the array: 2000000
cutoff: 510000      10 times Time :1051ms
cutoff: 520000      10 times Time :997ms
cutoff: 530000      10 times Time :1003ms
cutoff: 540000      10 times Time :983ms
cutoff: 550000      10 times Time :1053ms
cutoff: 560000      10 times Time :989ms
cutoff: 570000      10 times Time :974ms
cutoff: 580000      10 times Time :967ms
cutoff: 590000      10 times Time :983ms
cutoff: 600000      10 times Time :977ms
cutoff: 610000      10 times Time :987ms
cutoff: 620000      10 times Time :986ms
cutoff: 630000      10 times Time :970ms
cutoff: 640000      10 times Time :960ms
cutoff: 650000      10 times Time :957ms
cutoff: 660000      10 times Time :969ms
cutoff: 670000      10 times Time :958ms
cutoff: 680000      10 times Time :952ms
cutoff: 690000      10 times Time :961ms
cutoff: 700000      10 times Time :956ms
cutoff: 710000      10 times Time :973ms
cutoff: 720000      10 times Time :992ms
cutoff: 730000      10 times Time :965ms
cutoff: 740000      10 times Time :981ms
cutoff: 750000      10 times Time :957ms
cutoff: 760000      10 times Time :989ms
cutoff: 770000      10 times Time :1001ms
cutoff: 780000      10 times Time :1114ms
cutoff: 790000      10 times Time :1156ms
cutoff: 800000      10 times Time :1043ms
cutoff: 810000      10 times Time :997ms
cutoff: 820000      10 times Time :1073ms
cutoff: 830000      10 times Time :1008ms
cutoff: 840000      10 times Time :956ms
cutoff: 850000      10 times Time :951ms
cutoff: 860000      10 times Time :945ms
cutoff: 870000      10 times Time :970ms
cutoff: 880000      10 times Time :952ms
```

```
Degree of parallelism: 8
Size of the array: 4000000
cutoff: 510000      10 times Time :2165ms
cutoff: 520000      10 times Time :2048ms
cutoff: 530000      10 times Time :2138ms
cutoff: 540000      10 times Time :1822ms
cutoff: 550000      10 times Time :1747ms
cutoff: 560000      10 times Time :1967ms
cutoff: 570000      10 times Time :1880ms
cutoff: 580000      10 times Time :1722ms
cutoff: 590000      10 times Time :1759ms
cutoff: 600000      10 times Time :1835ms
cutoff: 610000      10 times Time :1806ms
cutoff: 620000      10 times Time :1654ms
cutoff: 630000      10 times Time :1801ms
cutoff: 640000      10 times Time :1738ms
cutoff: 650000      10 times Time :1653ms
cutoff: 660000      10 times Time :1736ms
cutoff: 670000      10 times Time :1744ms
cutoff: 680000      10 times Time :1707ms
cutoff: 690000      10 times Time :1772ms
cutoff: 700000      10 times Time :1784ms
cutoff: 710000      10 times Time :1790ms
cutoff: 720000      10 times Time :1785ms
cutoff: 730000      10 times Time :1813ms
cutoff: 740000      10 times Time :1853ms
cutoff: 750000      10 times Time :1753ms
cutoff: 760000      10 times Time :1728ms
cutoff: 770000      10 times Time :1764ms
cutoff: 780000      10 times Time :1754ms
cutoff: 790000      10 times Time :1787ms
cutoff: 800000      10 times Time :1725ms
cutoff: 810000      10 times Time :1764ms
cutoff: 820000      10 times Time :1732ms
cutoff: 830000      10 times Time :1718ms
cutoff: 840000      10 times Time :1742ms
cutoff: 850000      10 times Time :1929ms
cutoff: 860000      10 times Time :1747ms
cutoff: 870000      10 times Time :1771ms
cutoff: 880000      10 times Time :1778ms
```

```

Degree of parallelism: 16
Size of the array: 8000000
cutoff: 510000      10 times Time :3242ms
cutoff: 520000      10 times Time :3194ms
cutoff: 530000      10 times Time :3967ms
cutoff: 540000      10 times Time :4101ms
cutoff: 550000      10 times Time :3665ms
cutoff: 560000      10 times Time :3397ms
cutoff: 570000      10 times Time :3105ms
cutoff: 580000      10 times Time :3061ms
cutoff: 590000      10 times Time :3962ms
cutoff: 600000      10 times Time :3353ms
cutoff: 610000      10 times Time :3239ms
cutoff: 620000      10 times Time :4256ms
cutoff: 630000      10 times Time :3318ms
cutoff: 640000      10 times Time :3183ms
cutoff: 650000      10 times Time :3121ms
cutoff: 660000      10 times Time :3210ms
cutoff: 670000      10 times Time :3384ms
cutoff: 680000      10 times Time :3603ms
cutoff: 690000      10 times Time :3812ms
cutoff: 700000      10 times Time :3290ms
cutoff: 710000      10 times Time :3278ms
cutoff: 720000      10 times Time :3180ms
cutoff: 730000      10 times Time :3324ms
cutoff: 740000      10 times Time :3299ms
cutoff: 750000      10 times Time :3189ms
cutoff: 760000      10 times Time :3207ms
cutoff: 770000      10 times Time :3329ms
cutoff: 780000      10 times Time :3451ms
cutoff: 790000      10 times Time :3211ms
cutoff: 800000      10 times Time :3275ms
cutoff: 810000      10 times Time :3576ms
cutoff: 820000      10 times Time :3138ms
cutoff: 830000      10 times Time :3256ms
cutoff: 840000      10 times Time :3157ms
cutoff: 850000      10 times Time :3175ms
cutoff: 860000      10 times Time :3262ms
cutoff: 870000      10 times Time :3206ms

```

```

Degree of parallelism: 32
Size of the array: 16000000
cutoff: 510000      10 times Time :8378ms
cutoff: 520000      10 times Time :10597ms
cutoff: 530000      10 times Time :9460ms
cutoff: 540000      10 times Time :9215ms
cutoff: 550000      10 times Time :9186ms
cutoff: 560000      10 times Time :8942ms
cutoff: 570000      10 times Time :7792ms
cutoff: 580000      10 times Time :7267ms
cutoff: 590000      10 times Time :8033ms
cutoff: 600000      10 times Time :9385ms
cutoff: 610000      10 times Time :9681ms
cutoff: 620000      10 times Time :7650ms
cutoff: 630000      10 times Time :9481ms
cutoff: 640000      10 times Time :8507ms
cutoff: 650000      10 times Time :8509ms
cutoff: 660000      10 times Time :10009ms
cutoff: 670000      10 times Time :9419ms
cutoff: 680000      10 times Time :9340ms
cutoff: 690000      10 times Time :10793ms
cutoff: 700000      10 times Time :8907ms
cutoff: 710000      10 times Time :13059ms
cutoff: 720000      10 times Time :11173ms
cutoff: 730000      10 times Time :10775ms
cutoff: 740000      10 times Time :16045ms
cutoff: 750000      10 times Time :12111ms
cutoff: 760000      10 times Time :12978ms
cutoff: 770000      10 times Time :8502ms
cutoff: 780000      10 times Time :9338ms
cutoff: 790000      10 times Time :6297ms
cutoff: 800000      10 times Time :6969ms
cutoff: 810000      10 times Time :12023ms
cutoff: 820000      10 times Time :10202ms
cutoff: 830000      10 times Time :10107ms
cutoff: 840000      10 times Time :9766ms
cutoff: 850000      10 times Time :10977ms
cutoff: 860000      10 times Time :10797ms
cutoff: 870000      10 times Time :9145ms
cutoff: 880000      10 times Time :9672ms

```

2. Spreadsheet Data:

Here I have used spreadsheet data to monitor ratio of cutoff to array size and the degree of parallelism.

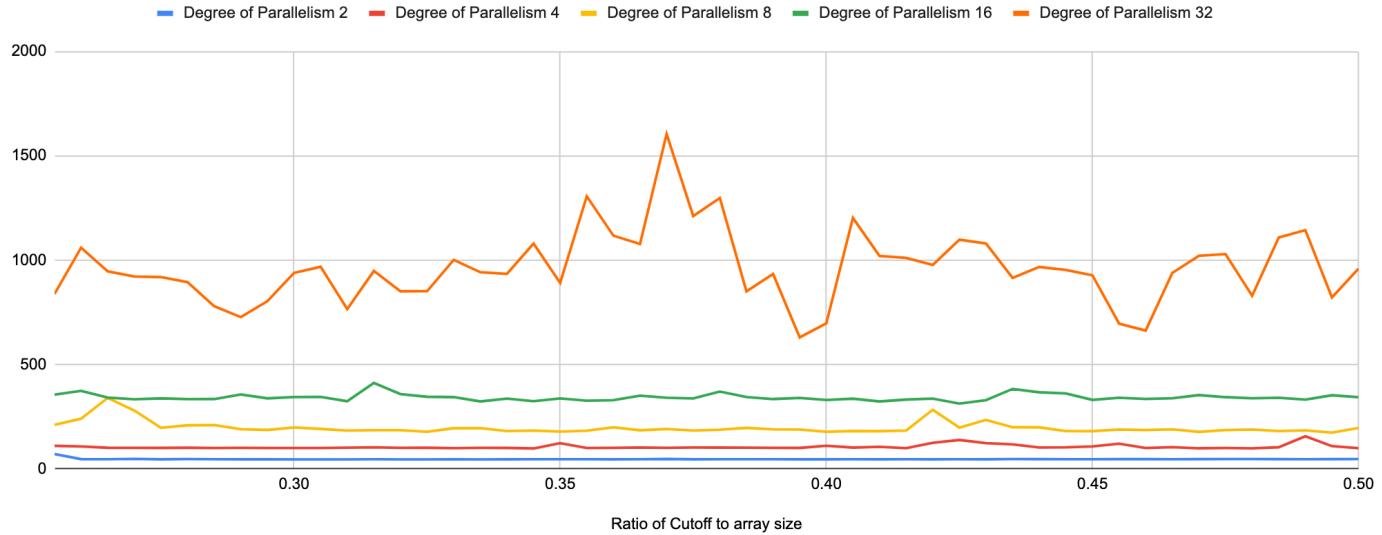
Ratio of Cutoff to array size	Degree of Parallelism 2	Degree of Parallelism 4	Degree of Parallelism 8	Degree of Parallelism 16	Degree of Parallelism 32
0.255	70.2	109.8	210.5	355.3	837.8
0.26	46.2	106.8	239.9	373.2	1059.7
0.265	45.6	100.8	339.6	341.5	946
0.27	47.2	100.6	279	333.1	921.5
0.275	45.3	99.9	196.3	337.5	918.6
0.28	46.7	100.7	208.4	333.9	894.2
0.285	45.5	99.6	209.3	334.1	779.2
0.29	45.3	100.2	189.9	355.9	726.7
0.295	45.3	99.5	186	337.5	803.3
0.3	44.8	99.8	198	343.6	938.5
0.305	45.2	99.3	191.1	344.3	968.1
0.31	45.1	101	183.1	323.8	765
0.315	45.9	102.6	184.7	411.4	948.1
0.32	45.1	100.4	184.6	357.8	850.7
0.325	45.2	100.7	177.3	345	850.9
0.33	45.4	99	194.3	343.4	1000.9
0.335	45.2	100.4	194.4	322.7	941.9
0.34	45.3	100	180.6	336.3	934
0.345	45.5	97.4	183	324.1	1079.3
0.35	45.7	122.7	178.2	336.8	890.7
0.355	45.5	99.5	182.4	326.5	1305.9
0.36	45.4	100.4	198.6	328.9	1117.3
0.365	45.5	101.9	184.4	350.5	1077.5
0.37	46.9	100.5	190.4	340.4	1604.5
0.375	45.3	101.9	183.3	336.9	1211.1
0.38	45.5	101.6	187.4	370.3	1297.8
0.385	45.6	101.1	195.9	343.7	850.2
0.39	45.7	100.2	188.8	334.1	933.8
0.395	45.4	100	187.9	339.3	629.7
0.4	45.4	110.1	177.5	329.9	696.9
0.405	45.7	101.7	181	335.8	1202.3
0.41	45.4	105.1	180.3	322.7	1020.2

0.415	45.6	98.6	183.4	331.7	1010.7
0.42	45.4	124.5	282.5	336.3	976.6
0.425	45.7	138	197.3	312.2	1097.7
0.43	45.4	122.8	234.1	329.1	1079.7
0.435	46.6	117.1	199.3	382.7	914.5
0.44	46.2	101.8	198.5	366.5	967.2
0.445	46	102.6	180.5	361.3	953
0.45	45.8	106.9	180.3	330.2	927.1
0.455	46.1	120.4	188.3	340.6	695.8
0.46	46.2	99.4	185.9	334.8	662.3
0.465	46	103.7	188.7	337.7	938.5
0.47	46.1	98.2	177	353.3	1020.6
0.475	46.6	99.4	185.5	343.2	1028.8
0.48	46.6	97.8	188.3	338	828.7
0.485	46.1	103.3	180.9	340.4	1108.4
0.49	45.9	155.8	184.1	331.6	1143.8
0.495	46.1	108.5	173.4	352.4	821.3
0.5	46.5	98.4	195.2	342.9	959.8

3. Graphical Representation:

The above data can be plotted on a graph :

Degree of Parallelism 2, Degree of Parallelism 4, Degree of Parallelism 8, Degree of Parallelism 16 and Degree of Parallelism 32



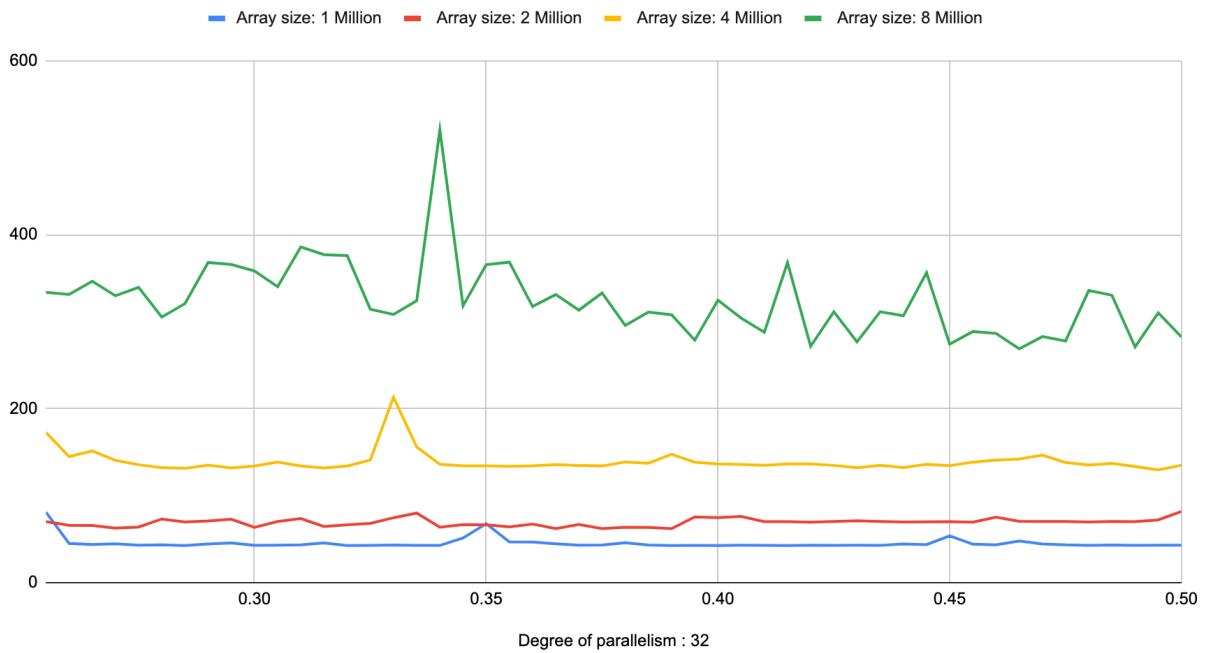
4. Evidence to study the behavior of the algorithm when thread count is 32 and size of the array increases in each iteration.

For this I have created a new Test32DegreeParallelism.

Ratio of Cutoff to array size	Array size: 1 Million		Array size: 2 Million	Array size: 4 Million	Array size: 8 Million
	0.255	81	70	172.7	333.8
0.26	45	65.9	144.8	331.4	
0.265	43.8	65.6	151.3	346.5	
0.27	44.6	62.7	140.6	329.8	
0.275	42.9	63.8	135.5	339.5	
0.28	43.4	73.1	132.1	305.3	
0.285	42.6	69.7	131.3	320.8	
0.29	44.3	70.8	134.9	368.1	
0.295	45.6	72.8	131.8	365.8	
0.3	42.7	63.5	134	358.4	
0.305	43	70.2	138.5	340.3	
0.31	43.3	73.6	134	386	
0.315	45.5	64.4	131.8	377.1	

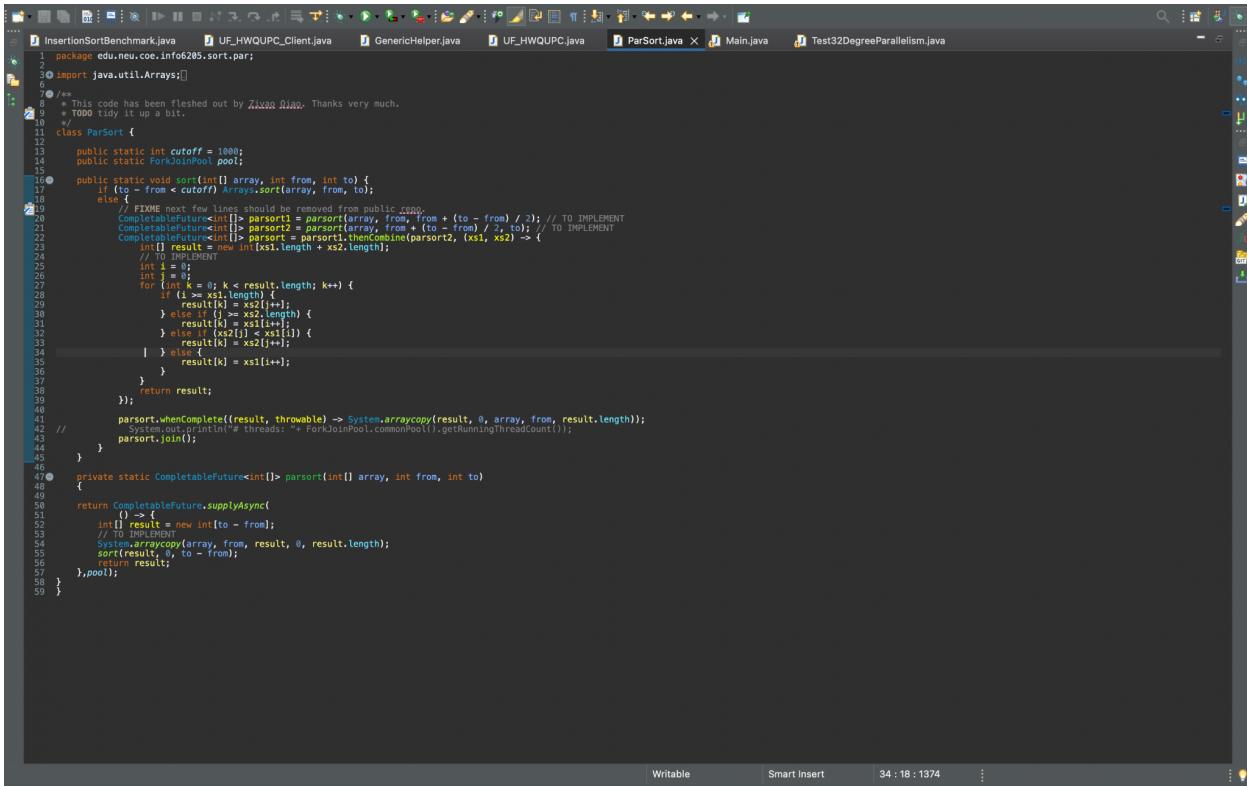
0.32	42.5	66.5	134	376
0.325	42.7	68.1	140.9	314.3
0.33	43.1	74.5	213.1	308.3
0.335	42.7	79.9	155.9	324.1
0.34	42.7	63.6	136	520.5
0.345	51.2	66.7	134.2	318.2
0.35	67.7	66.5	134.2	365.6
0.355	46.8	64.1	133.5	368.5
0.36	46.5	67.3	134.2	317.4
0.365	44.5	62	135.6	331.2
0.37	42.9	66.7	134.6	313.3
0.375	43.2	62	134.1	333.1
0.38	45.7	63.6	138.7	295.7
0.385	43.2	63.5	137.3	311
0.39	42.6	62	147.6	307.9
0.395	42.8	75.4	138.4	278.8
0.4	42.6	74.6	136.5	324.7
0.405	42.9	76.1	135.9	304.1
0.41	42.8	70	134.9	287.8
0.415	42.6	70	136.4	367.9
0.42	43	69.5	136.4	271.7
0.425	42.8	70.3	134.7	311.4
0.43	43	71.1	132.1	276.8
0.435	42.8	70.2	134.8	311.4
0.44	44.4	69.6	132.2	306.8
0.445	43.6	69.8	135.9	356.3
0.45	53.7	70	134.5	274.1
0.455	44.1	69.5	138.5	288.7
0.46	43.3	75.3	140.8	286.5
0.465	47.7	70.4	142.1	268.8
0.47	44.4	70.2	146.6	282.9
0.475	43.3	70.3	138	277.8
0.48	42.8	69.7	135.3	335.9
0.485	43.2	70.3	137	330.4
0.49	42.7	70.1	133.4	271
0.495	43	72.1	129.6	310.2
0.5	42.9	81.7	134.9	282.4

1 Million, 2 Million, 4 Million and 8 Million



Code Snippet:

The following are the changes made to ParSort class



A screenshot of an IDE showing the code for the ParSort class. The code implements a parallel sorting algorithm using Java's ForkJoinPool. It includes methods for sorting arrays and returning CompletableFuture<int[]>. The code is annotated with TODO and FIXME comments, indicating areas for improvement.

```
1 package edu.neu.coe.info6205.sort.par;
2
3 import java.util.Arrays;
4
5 /**
6 * This code has been fleshed out by Zixuan Qiao. Thanks very much.
7 * + TODO tidy it up a bit.
8 */
9
10 class ParSort {
11
12     public static int cutoff = 1000;
13
14     public static ForkJoinPool pool;
15
16     public static void sort(int[] array, int from, int to) {
17         if (to - from < cutoff) Arrays.sort(array, from, to);
18         else {
19             // FIXME next few lines should be removed from public code.
20             CompletableFuture<int[]> parsort = parsort(array, from, from + (to - from) / 2);
21             CompletableFuture<int[]> parsort2 = parsort.thenApply((array, from + (to - from) / 2, to) -> {
22                 int result = new int[xs1.length + xs2.length];
23                 // TO IMPLEMENT
24                 int i = 0;
25                 int j = 0;
26                 for (int k = 0; k < result.length; k++) {
27                     if (i >= xs1.length)
28                         result[k] = xs2[j++];
29                     else if (j >= xs2.length)
30                         result[k] = xs1[i++];
31                     else if (xs2[j] < xs1[i]) {
32                         result[k] = xs2[j++];
33                     } else {
34                         result[k] = xs1[i++];
35                     }
36                 }
37             });
38             return result;
39         };
40         parsort.whenComplete((result, throwable) -> System.arraycopy(result, 0, array, from, result.length));
41         // System.out.println("threads: " + ForkJoinPool.commonPool().getRunningThreadCount());
42         pool.submit(parsort);
43     }
44
45     private static CompletableFuture<int[]> parsort(int[] array, int from, int to) {
46
47         return CompletableFuture.supplyAsync(
48             () -> {
49                 int[] result = new int[to - from];
50                 // TO IMPLEMENT
51                 System.arraycopy(array, from, result, 0, result.length);
52                 sort(result, 0, to - from);
53                 return result;
54             }, pool);
55     }
56 }
57
58 }
```

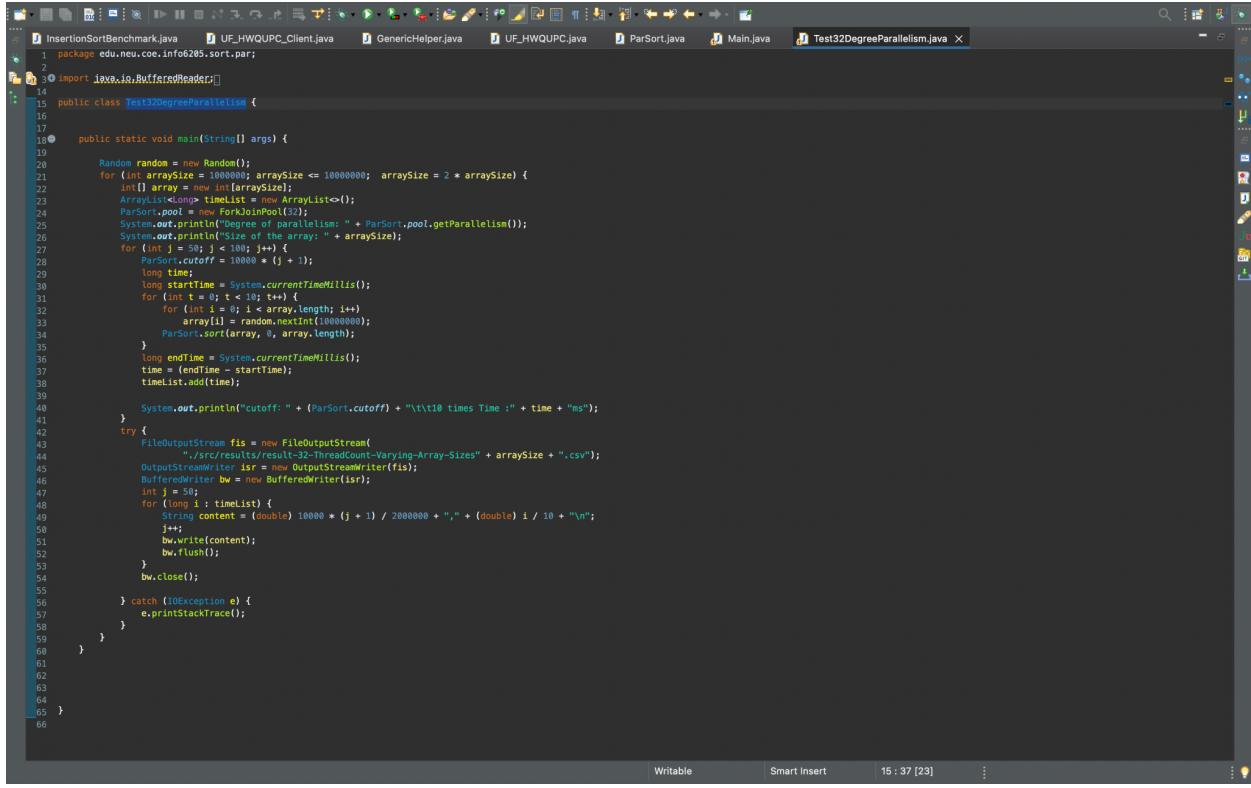
The following are the changes made to Main class

```

28     public static void main(String[] args) {
29         processArgs(args);
30         BufferedReader reader = new BufferedReader(new InputStreamReader(System.in));
31         int initialThreadCount = 0;
32         int initialArraySize = 0; // 1000000
33         try {
34             System.out.println("Enter initial thread count :");
35             String thread = reader.readLine();
36             initialThreadCount = Integer.parseInt(thread);
37             System.out.println("Enter initial array size : ");
38             String size = reader.readLine();
39             initialArraySize = Integer.parseInt(size);
40         } catch (IOException e) {
41             e.printStackTrace();
42         }
43         Random random = new Random();
44         int arraySize = initialArraySize;
45         for (int threadCount = initialThreadCount; threadCount <= 32; threadCount = 2 * threadCount, arraySize = 2
46             * arraySize) {
47             int[] array = new int[arraySize];
48             ArrayList<long> timeList = new ArrayList<>();
49             ForkJoinPool pool = new ForkJoinPool(threadCount);
50             System.out.println("Degree of parallelism: " + pool.getParallelism());
51             System.out.println("Size of the array: " + arraySize);
52             for (int j = 50; j < 100; j++) {
53                 ParSort.cutoff = 10000 * (j + 1);
54                 long startTime = System.currentTimeMillis();
55                 for (int i = 0; i < array.length; i++) {
56                     array[i] = random.nextInt(1000000);
57                 }
58                 long endTime = System.currentTimeMillis();
59                 time = (endTime - startTime);
60                 timeList.add(time);
61             }
62             System.out.println("cutoff: " + (ParSort.cutoff) + "\t\t" + "Time :" + time + "ms");
63         }
64         try {
65             FileOutputStream fis = new FileOutputStream(
66                 "./src/results/result-" + threadCount + "-threads=" + arraySize + ".csv");
67             OutputStreamWriter isr = new OutputStreamWriter(fis);
68             BufferedWriter bw = new BufferedWriter(isr);
69             int j = 50;
70             for (long i : timeList) {
71                 String content = (double) 10000 * (j + 1) / 2000000 + "," + (double) i / 10 + "\n";
72                 j++;
73                 bw.write(content);
74                 bw.flush();
75             }
76             bw.close();
77         } catch (IOException e) {
78             e.printStackTrace();
79         }
80     }

```

I also created a new test class for my own experimentation:



The screenshot shows an IDE interface with multiple tabs open. The active tab is 'Test32DegreeParallelism.java'. The code in this file is as follows:

```
1 package edu.neu.coe.info6205.sort.par;
2
3 import java.io.BufferedReader;
4
5 public class Test32DegreeParallelism {
6
7     public static void main(String[] args) {
8
9         Random random = new Random();
10        for (int arraySize = 1000000; arraySize <= 10000000; arraySize = 2 * arraySize) {
11            int[] array = new int[arraySize];
12            ArrayList<Long> timeList = new ArrayList<Long>();
13            ParSort.pool = new ForkJoinPool(32);
14            System.out.println("Degree of parallelism: " + ParSort.pool.getParallelism());
15            System.out.println("Size of the array: " + arraySize);
16            for (int j = 50; j < 100; j++) {
17                ParSort.cutoff = 1000 * (j + 1);
18                long startTime = System.currentTimeMillis();
19                for (int t = 0; t < 10; t++) {
20                    for (int i = 0; i < array.length; i++) {
21                        array[i] = random.nextInt(10000000);
22                    }
23                    ParSort.sort(array, 0, array.length);
24                }
25                long endTime = System.currentTimeMillis();
26                long time = (endTime - startTime);
27                timeList.add(time);
28            }
29            System.out.println("cutoff: " + (ParSort.cutoff) + "\t\t10 times Time :" + time + "ms");
30        }
31        try {
32            FileOutputStream file = new FileOutputStream(
33                "/src/results/result-32-ThreadCount-Varying-Array-Sizes" + arraySize + ".csv");
34            OutputStreamWriter isr = new OutputStreamWriter(file);
35            BufferedWriter bw = new BufferedWriter(isr);
36            int j = 50;
37            for (long i : timeList) {
38                String content = (double) 10000 * (j + 1) / 2800000 + "," + (double) i / 10 + "\n";
39                j++;
40                bw.write(content);
41                bw.flush();
42            }
43            bw.close();
44        } catch (IOException e) {
45            e.printStackTrace();
46        }
47    }
48}
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66}
```